

Lab 1

Sam Burch

2023-02-04

Question 1

```
library(Lahman)
```

```
## Warning: package 'Lahman' was built under R version 4.2.2
```

```
library(broom)
```

```
## Warning: package 'broom' was built under R version 4.2.2
```

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.3.6      v purrr   0.3.4
## v tibble  3.1.8      v dplyr  1.0.9
## v tidyr   1.2.0      v stringr 1.4.1
## v readr   2.1.2      v forcats 0.5.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
Teams = as_tibble(Teams)
```

```
df_1 = Teams |>
  select(yearID, W, L, G, AB, H, X2B, X3B, HR, BB, HBP, SF, HA, HRA,
         BBA, SOA, IPouts, FP, R, RA) |>
  filter(yearID >= 1900) |>
  replace_na(list(HBP = 0, SF = 0)) |>
  mutate(RD = (R - RA) / G,
         X1B = H - (X2B + X3B + HR)) |>
  mutate(OBP = (H + BB + HBP) / (AB + BB + HBP + SF)) |>
  mutate(SLG = (X1B + 2*X2B + 3*X3B + 4*HR) / AB) |>
  mutate(OPS = OBP + SLG) |>
  mutate(IP = IPouts / 3) |>
  mutate(WHIP = (BBA + HA) / IP) |>
  mutate(FIP = (13*HRA + 3*BBA - 2*SOA) / IP)
df_1
```

```
## # A tibble: 2,610 x 28
##   yearID      W      L      G      AB      H      X2B      X3B      HR      BB      HBP      SF
##   <int> <int> <int> <int> <int> <int> <int> <int> <int> <int> <int> <int>
## 1  1900     82     54    141   4860   1423    199     81     26    421     81     0
## 2  1900     66     72    142   4952   1403    163     68     48    395     45     0
## 3  1900     65     75    146   4907   1276    202     51     33    343     65     0
## 4  1900     62     77    144   5026   1335    178     83     33    333     50     0
## 5  1900     60     78    141   4724   1317    177     61     23    369     56     0
## 6  1900     75     63    141   4969   1439    187     82     29    440     72     0
## 7  1900     79     60    140   4817   1312    185    100     26    327     63     0
## 8  1900     65     75    141   4877   1420    141     81     36    406     81     0
## 9  1901     68     65    134   4589   1348    179    111     24    369     52     0
## 10 1901     79     57    138   4866   1353    183    104     37    331     47     0
## # ... with 2,600 more rows, and 16 more variables: HA <int>, HRA <int>,
## #   BBA <int>, SOA <int>, IPouts <int>, FP <dbl>, R <int>, RA <int>, RD <dbl>,
## #   X1B <int>, OBP <dbl>, SLG <dbl>, OPS <dbl>, IP <dbl>, WHIP <dbl>, FIP <dbl>
```

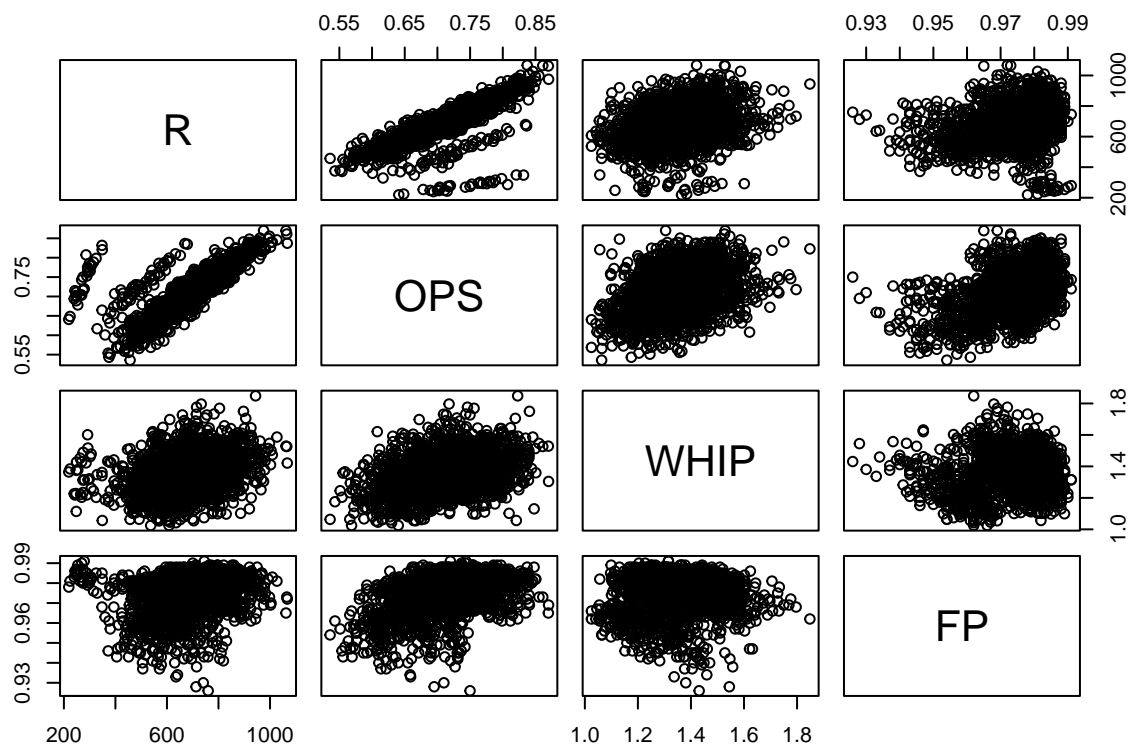
```
mlr_1 = lm(R ~ OPS + WHIP + FP, data = df_1)
summary(mlr_1)
```

```
##
## Call:
## lm(formula = R ~ OPS + WHIP + FP, data = df_1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -563.74  -13.16    9.62   29.36  133.98
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1187.52     150.14   7.909 3.79e-15 ***
## OPS          1958.31      29.83  65.659 < 2e-16 ***
## WHIP         -45.19      12.32  -3.668 0.00025 ***
## FP          -1882.77     159.93 -11.772 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 66.43 on 2606 degrees of freedom
## Multiple R-squared:  0.6697, Adjusted R-squared:  0.6693
## F-statistic: 1761 on 3 and 2606 DF, p-value: < 2.2e-16
```

R² here is ok, but let us now test for diagnostics.

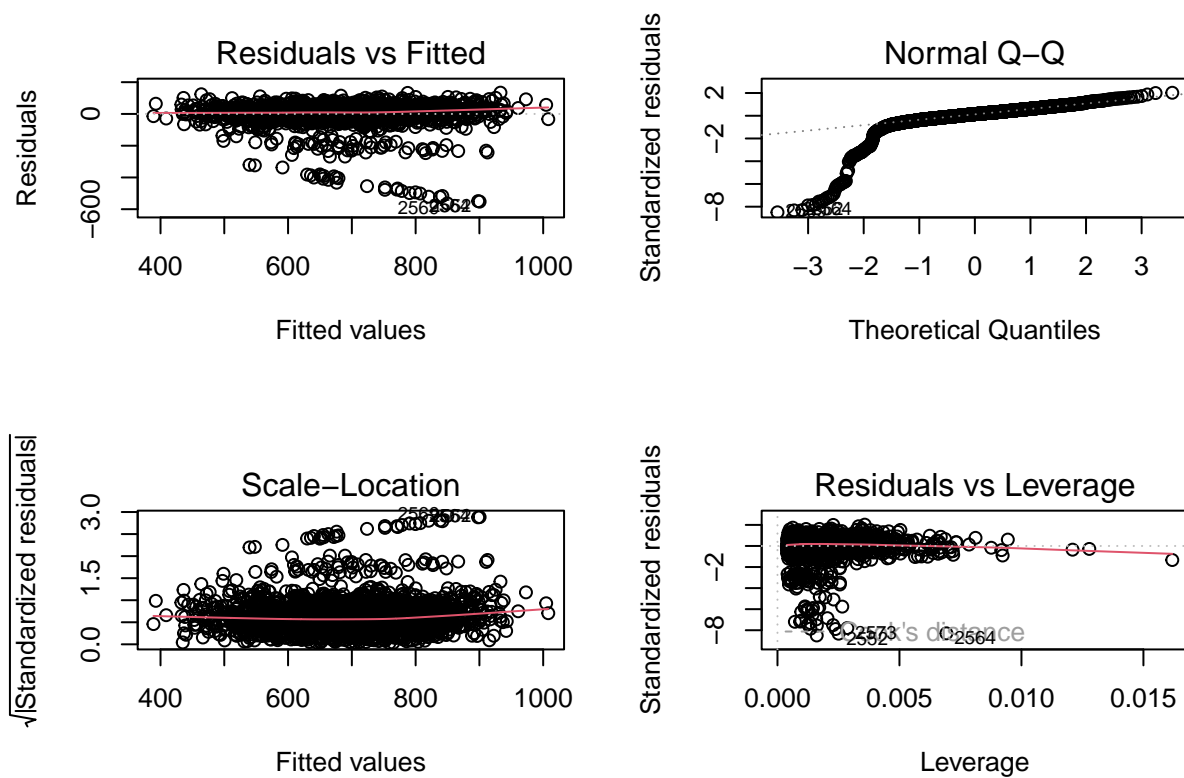
Diagnostics

```
pairs(df_1 |> dplyr::select(R, OPS, WHIP, FP))
```



The linearity assumption holds, as there is clearly a linear relationship between runs and each predictor in the model.


```
par(mfrow = c(2, 2))
plot(mlr_1)
```



As we check that the errors are normally distributed with constant variance and mean 0, only mean 0 seems to suffice. There seems to be deviations from normality and constant variance assumptions. Also, a large right tail in residuals.

Problematic Residuals

```
par(mfrow = c(1, 1))
qqnorm(residuals(mlr_1))
qqline(residuals(mlr_1))
abline(h = -50, lty = 2, col = "red")
```



It seems like certain residuals are causing the normality and constant variance assumptions to fail.

```
df_1 = augment(mlr_1, df_1)
```

```
df_1 |>
  filter(.resid < -100) |>
  group_by(yearID) |>
  summarise(n())
```

```
## # A tibble: 6 x 2
##   yearID 'n()'
##   <int> <int>
## 1  1918     5
## 2  1919     3
## 3  1921     1
## 4  1981    26
## 5  1994    28
## 6  2020    30
```

```
df_1 |>
  filter(.resid < -50) |>
  group_by(yearID) |>
  summarise(n = n()) |>
  filter(n >= 10)
```

```
## # A tibble: 6 x 2
```

```
##   yearID      n
##   <int> <int>
## 1   1918     16
## 2   1919     13
## 3   1981     26
## 4   1994     28
## 5   1995     18
## 6   2020     30
```

These years are where the issues are coming from. Let us try to fix this problem by adding in categorical variables for these years.

Fixed Model

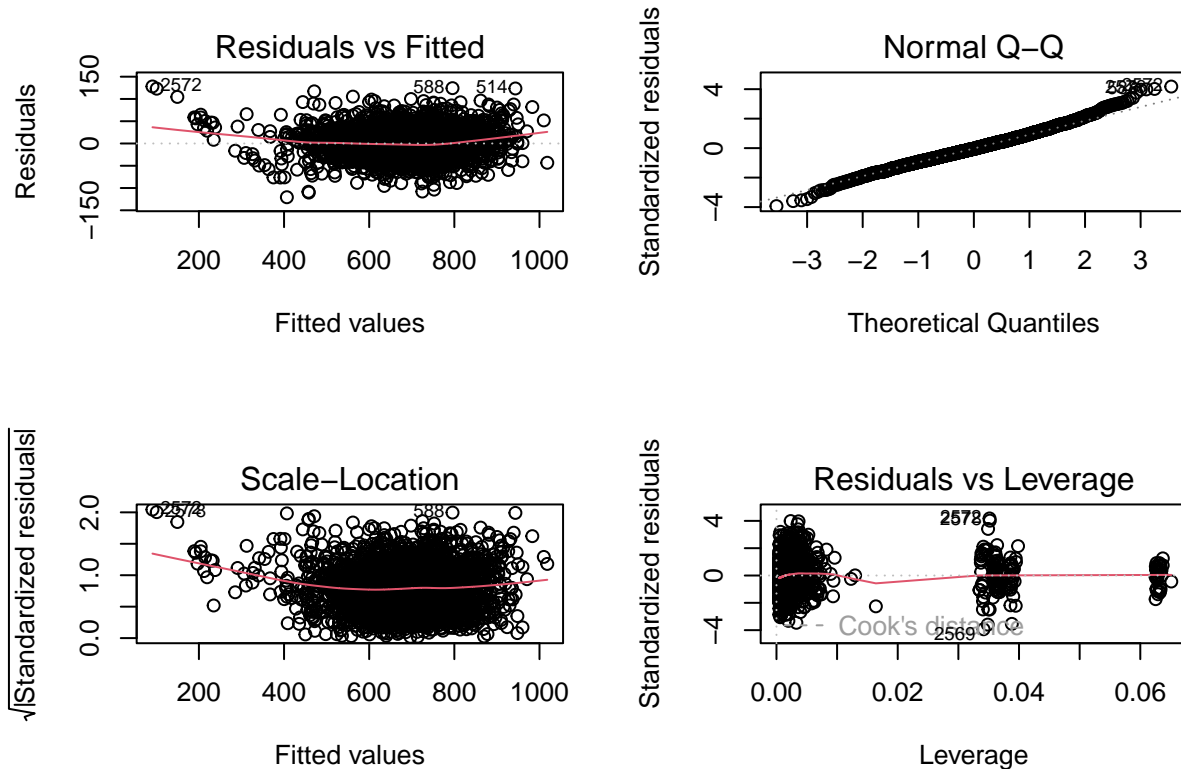
```
df_1 = df_1 |>
  mutate(year_1918 = ifelse(yearID == 1918, 1, 0),
         year_1919 = ifelse(yearID == 1919, 1, 0),
         year_1981 = ifelse(yearID == 1981, 1, 0),
         year_1994 = ifelse(yearID == 1994, 1, 0),
         year_1995 = ifelse(yearID == 1995, 1, 0),
         year_2020 = ifelse(yearID == 2020, 1, 0))

mlr_2 = lm(R ~ OPS + WHIP + FP + year_1918 + year_1919 + year_1981 + year_1994 + year_1995 + year_2020,
           summary(mlr_2))
```

```
##
## Call:
## lm(formula = R ~ OPS + WHIP + FP + year_1918 + year_1919 + year_1981 +
##     year_1994 + year_1995 + year_2020, data = df_1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -120.670  -20.487   -1.574    18.985   128.068
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    625.787     71.396   8.765 <2e-16 ***
## OPS            1970.837    14.110 139.673 <2e-16 ***
## WHIP            -56.189     5.818  -9.658 <2e-16 ***
## FP           -1288.836    75.875 -16.986 <2e-16 ***
## year_1918      -108.151     7.904 -13.683 <2e-16 ***
## year_1919       -79.908     7.872 -10.151 <2e-16 ***
## year_1981      -218.064     6.176 -35.309 <2e-16 ***
## year_1994      -222.023     5.967 -37.206 <2e-16 ***
## year_1995       -70.798     5.960 -11.878 <2e-16 ***
## year_2020      -460.827     5.762 -79.973 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 31.22 on 2600 degrees of freedom
## Multiple R-squared:  0.9272, Adjusted R-squared:  0.927
## F-statistic: 3681 on 9 and 2600 DF, p-value: < 2.2e-16
```

Here, the R^2 is much higher, which is a good sign!

```
par(mfrow = c(2, 2))
plot(mlr_2)
```



All assumptions met now. Thus, this is a much better model than before.

What went wrong is these outlier seasons, due to different circumstances, had drastically different predictor values related to runs. So, after adding in the categorical variables for such years, this problem was solved!

Principled Rescaling / New Model

```
teams_list = Teams |>
  select(yearID, W, L, G, AB, H, X2B, X3B, HR, BB, HBP, SF, HA, HRA,
         BBA, SOA, IPouts, FP, R, RA) |>
  replace_na(list(HBP = 0, SF = 0)) |>
  mutate(RD = (R - RA) / G) |>
  mutate(X1B = H - (X2B + X3B + HR)) |>
  mutate(OBP = (H + BB + HBP) / (AB + BB + HBP + SF)) |>
  mutate(SLG = (X1B + 2*X2B + 3*X3B + 4*HR) / AB) |>
  mutate(OPS = OBP + SLG) |>
  mutate(IP = IPouts / 3) |>
  mutate(WHIP = (BBA + HA) / IP) |>
  split(Teams$yearID) |>
  lapply(avgOBP = (sum(H) + sum(BB) + sum(HBP)) / (sum(AB) + sum(BB) + sum(HBP) + sum(SF)),
        avgSLG = (sum(X1B) + 2*sum(X2B) + 3*sum(X3B) + 4*sum(HR)) / sum(AB),
        avgOPS = avgOBP + avgSLG,
```

```

    avgWHIP = (sum(BBA) + sum(HA)) / sum(IP),
    #avgFP here is only rough approximation here
    avgFP = mean(FP),
    mutate) |>
lapply(OPSscale = OPS / avgOPS,
       WHIPscale = avgWHIP / WHIP,
       FPscale = avgFP / FP,
       mutate)

df_2 = do.call('rbind', teams_list)

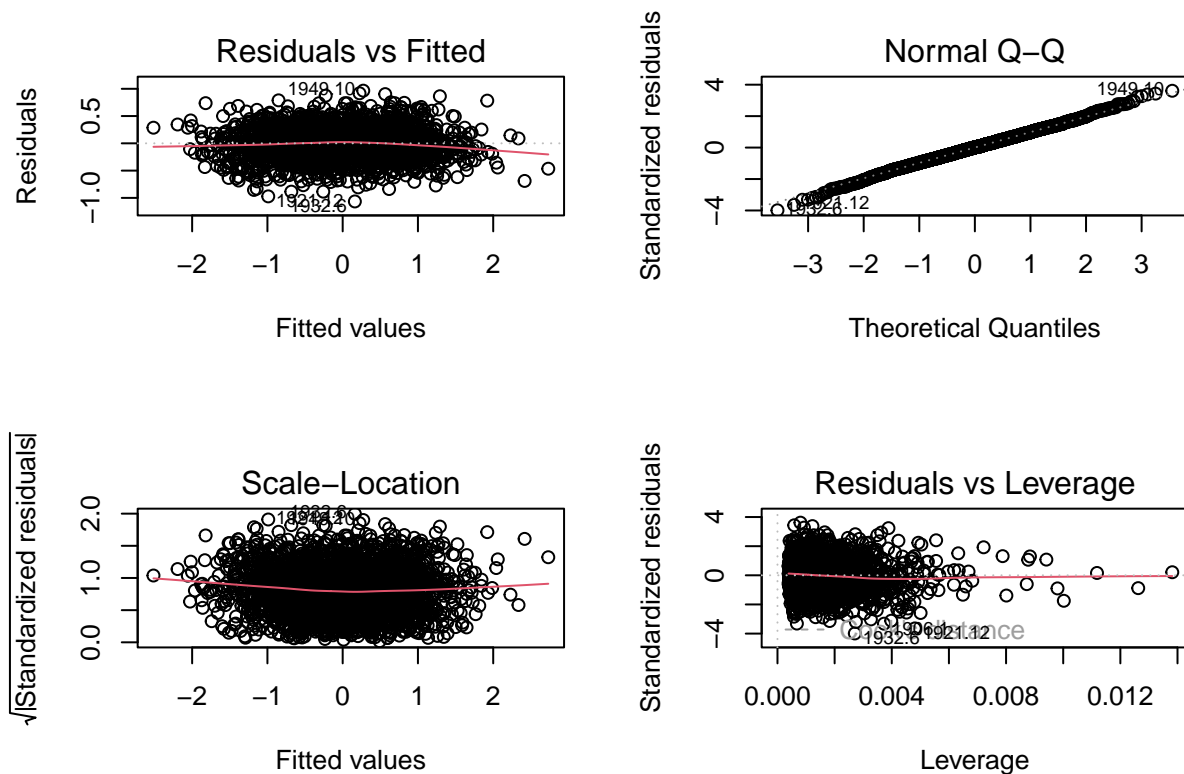
df_2 = df_2 |>
  filter(yearID >= 1900)

mlr_3 = lm(RD ~ OPSscale + WHIPscale + FPscale, data = df_2)
summary(mlr_3)

##
## Call:
## lm(formula = RD ~ OPSscale + WHIPscale + FPscale, data = df_2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.06511 -0.17520  0.00438  0.17697  0.96536
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   6.62214     1.67890   3.944 8.21e-05 ***
## OPSscale      8.99139     0.10691  84.104 < 2e-16 ***
## WHIPscale     7.02806     0.08821  79.677 < 2e-16 ***
## FPscale     -22.66376     1.61910 -13.998 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2673 on 2606 degrees of freedom
## Multiple R-squared:  0.8754, Adjusted R-squared:  0.8752
## F-statistic: 6102 on 3 and 2606 DF, p-value: < 2.2e-16

par(mfrow = c(2, 2))
plot(mlr_3)

```

The rescaling allowed for a better model compared to the one in the notes ($RD \sim OPS + WHIP + FP$) because the R^2 is higher and all the model assumptions are met. This rescaling allowed the predictors to be accounted for their specific years. Meaning, if a year is different from the overall average, this accounts for that difference.

Question 2

Batting

```
Batting = as_tibble(Batting)
People = as_tibble(People)

batters_3 = People |>
  filter(playerID == 'troutmi01' |
         playerID == 'goldspa01' |
         playerID == 'justida01') |>
  dplyr::select(playerID, nameFirst, nameLast) |>
  pull(playerID)

bat_seasons = Batting |>
  filter(playerID %in% batters_3) |>
  group_by(playerID, yearID) |>
  mutate(X1B = H - (X2B + X3B + HR)) |>
  mutate(OBP = (H + BB + HBP) / (AB + BB + HBP + SF)) |>
  mutate(SLG = (X1B + 2*X2B + 3*X3B + 4*HR) / AB) |>
  mutate(OPS = OBP + SLG) |>
  mutate(SB_pct = (SB / (SB + CS))) |>
  mutate(season = case_when((playerID == 'justida01' & yearID == 1989) ~ 1,
```

```

        (playerID == 'justida01' & yearID != 1989) ~ (yearID - 1988),
        (playerID == 'goldspa01' & yearID == 2011) ~ 1,
        (playerID == 'goldspa01' & yearID != 2011) ~ (yearID - 2010),
        (playerID == 'troutmi01' & yearID == 2011) ~ 1,
        (playerID == 'troutmi01' & yearID != 2011) ~ (yearID - 2010),
    )) |>
    dplyr::select(playerID, yearID, season, G:SO, SB_pct, OBP, SLG, OPS)
bat_seasons

```

```

## # A tibble: 37 x 19
## # Groups:   playerID, yearID [36]
##   playerID yearID season      G    AB     R     H   X2B   X3B    HR   RBI    SB
##   <chr>      <int>  <dbl> <int> <int> <int> <int> <int> <int> <int> <int> <int>
## 1 justida01  1989      1    16    51     7    12     3     0     1     3     2
## 2 justida01  1990     2   127   439    76   124    23     2    28    78    11
## 3 justida01  1991     3   109   396    67   109    25     1    21    87     8
## 4 justida01  1992     4   144   484    78   124    19     5    21    72     2
## 5 justida01  1993     5   157   585    90   158    15     4    40   120     3
## 6 justida01  1994     6   104   352    61   110    16     2    19    59     2
## 7 justida01  1995     7   120   411    73   104    17     2    24    78     4
## 8 justida01  1996     8    40   140    23    45     9     0     6    25     1
## 9 justida01  1997     9   139   495    84   163    31     1    33   101     3
## 10 justida01 1998    10   146   540    94   151    39     2    21    88     9
## # ... with 27 more rows, and 7 more variables: CS <int>, BB <int>, SO <int>,
## #   SB_pct <dbl>, OBP <dbl>, SLG <dbl>, OPS <dbl>

```

```


Batting |>
  filter(playerID %in% batters_3) |>
  mutate(X1B = H - (X2B + X3B + HR)) |>
  group_by(playerID) |>
  summarise(seasons = n(),
            G = sum(G),
            AB = sum(AB),
            R = sum(R),
            H = sum(H),
            X2B = sum(X2B),
            X3B = sum(X3B),
            HR = sum(HR),
            RBI = sum(RBI),
            SB = sum(SB),
            CS = sum(CS),
            BB = sum(BB),
            SO = sum(SO),

            X1B = sum(X1B),
            HBP = sum(HBP),
            SF = sum(SF)
  ) |>
  mutate(OBP = (H + BB + HBP) / (AB + BB + HBP + SF)) |>
  mutate(SLG = (X1B + 2*X2B + 3*X3B + 4*HR) / AB) |>
  mutate(OPS = OBP + SLG) |>
  mutate(SB_pct = (SB / (SB + CS))) |>
  dplyr::select(playerID, seasons, G:SO, OBP:SB_pct)

```

```
## # A tibble: 3 x 18
##   playerID seasons      G      AB      R      H     X2B     X3B     HR     RBI     SB     CS
##   <chr>      <int> <int> <int> <int> <int> <int> <int> <int> <int> <int> <int>
## 1 goldspa01     11  1469  5366   939  1572   341    22   280   927   140   33
## 2 justida01     15  1610  5625   929  1571   280    24   305  1017    53   46
## 3 troutmi01     11  1288  4656   967  1419   268    49   310   816   203   37
## # ... with 6 more variables: BB <int>, SO <int>, OBP <dbl>, SLG <dbl>,
## #   OPS <dbl>, SB_pct <dbl>
```

```
ggplot(bat_seasons, aes(x = season, y = OPS)) +
  labs(
    x = 'Seasons',
    y = 'OPS',
    title = "OPS Career Trajectories",
    subtitle = 'Paul Goldschmidt, David Justice, and Mike Trout',
    caption = 'By: Sam Burch | Data @Lahman'
  ) +
  geom_point(aes(size = AB, color = playerID, alpha = .8)) +
  geom_line(aes(color = playerID, alpha = .8), data = bat_seasons |> filter(playerID == 'justida01')) +
  geom_line(aes(color = playerID, alpha = .8), data = bat_seasons |> filter(playerID == 'goldspa01')) +
  geom_line(aes(color = playerID, alpha = .8), data = bat_seasons |> filter(playerID == 'troutmi01')) +
  geom_hline(yintercept = .729, color = 'darkgrey', linetype = 2)
```



Note: The dotted grey line represents the approximate (modern) average OPS in MLB – .729.

We can see just how great all three batters have been! Basically every year these batters have been in the league, they've had a great OPS – well above average. Trout's even been hovering around 1.1 OPS the last

few years, which is absurd. For Justice, this graph shows the fall off (as he got older). Lastly, it only took till their second year for all three to be very good hitters.

Pitching

```
Pitching = as_tibble(Pitching)

pitchers_3 = People |>
  filter(playerID == 'bumgama01' |
         playerID == 'kershcl01' |
         playerID == 'riverma01') |>
  dplyr::select(playerID, nameFirst, nameLast) |>
  pull(playerID)

pitch_seasons = Pitching |>
  filter(playerID %in% pitchers_3) |>
  group_by(playerID, yearID) |>
  mutate(IP = IPouts / 3) |>
  mutate(WHIP = (BB + H) / IP) |>
  mutate(SO_per_9 = SO / (IP * 9)) |>
  mutate(SO_per_BB = SO / BB) |>
  mutate(season = case_when((playerID == 'riverma01' & yearID == 1995) ~ 1,
                           (playerID == 'riverma01' & yearID != 1995) ~ (yearID - 1994),
                           (playerID == 'kershcl01' & yearID == 2008) ~ 1,
                           (playerID == 'kershcl01' & yearID != 2008) ~ (yearID - 2007),
                           (playerID == 'bumgama01' & yearID == 2009) ~ 1,
                           (playerID == 'bumgama01' & yearID != 2009) ~ (yearID - 2008),
                           )) |>
  dplyr::select(playerID:yearID, season, W:L, IPouts:SO, HBP, ERA, WHIP,
               SO_per_9, SO_per_BB, IP)
pitch_seasons
```

```
## # A tibble: 46 x 17
## # Groups:   playerID, yearID [46]
##   playerID yearID season    W    L IPouts    H    ER    HR    BB    SO    HBP
##   <chr>      <int> <dbl> <int> <int> <int> <int> <int> <int> <int> <int> <int>
## 1 riverma~  1995     1     5     3   201    71    41    11    30    51     2
## 2 riverma~  1996     2     8     3   323    73    25     1    34   130     2
## 3 riverma~  1997     3     6     4   215    65    15     5    20    68     0
## 4 riverma~  1998     4     3     0   184    48    13     3    17    36     1
## 5 riverma~  1999     5     4     3   207    43    14     2    18    52     3
## 6 riverma~  2000     6     7     4   227    58    24     4    25    58     0
## 7 riverma~  2001     7     4     6   242    61    21     5    12    83     1
## 8 riverma~  2002     8     1     4   138    35    14     3    11    41     2
## 9 riverma~  2003     9     5     2   212    61    13     3    10    63     4
## 10 riverma~ 2004    10     4     2   236    65    17     3    20    66     5
## # ... with 36 more rows, and 5 more variables: ERA <dbl>, WHIP <dbl>,
## #   SO_per_9 <dbl>, SO_per_BB <dbl>, IP <dbl>
```

```
Pitching |>
  filter(playerID %in% pitchers_3) |>
  mutate(IP = IPouts / 3) |>
```

```

group_by(playerID) |>
summarise(
  seasons = n(),
  W = sum(W),
  L = sum(L),
  IPouts = sum(IPouts),
  H = sum(H),
  ER = sum(ER),
  HR = sum(HR),
  BB = sum(BB),
  HBP = sum(HBP),
  SO = sum(SO),

  IP = sum(IP)
) |>
mutate(ERA = (9 * ER) / IP) |>
mutate(WHIP = (BB + H) / IP) |>
mutate(SO_per_9 = SO / (IP * 9)) |>
mutate(SO_per_BB = SO / BB) |>
dplyr::select(-IP)

```

```

## # A tibble: 3 x 15
##   playerID seasons      W      L IPouts      H      ER      HR      BB      HBP      SO      ERA
##   <chr>      <int> <int> <int> <int> <int> <int> <int> <int> <int> <int> <dbl>
## 1 bumgama01     13    127    106   6102   1803    748    229    480     79   1948   3.31
## 2 kershcl01     14    185     84   7364   1859    679    196    606     37   2670   2.49
## 3 riverma01     19     82     60   3851    998    315     71    286     46   1173   2.21
## # ... with 3 more variables: WHIP <dbl>, SO_per_9 <dbl>, SO_per_BB <dbl>

```

```

People |>
  filter(playerID == 'bumgama01' |
         playerID == 'kershcl01' |
         playerID == 'riverma01') |>
  dplyr::select(playerID, nameFirst, nameLast)

```

```

## # A tibble: 3 x 3
##   playerID nameFirst nameLast
##   <chr>      <chr>      <chr>
## 1 bumgama01 Madison    Bumgarner
## 2 kershcl01 Clayton    Kershaw
## 3 riverma01 Mariano     Rivera

```

```


ggplot(pitch_seasons, aes(season, WHIP)) +
  labs(
    x = 'Seasons',
    y = 'WHIP',
    title = "WHIP Career Trajectories",
    subtitle = 'Madison Bumgarner, Clayton Kershaw, and Mariano Rivera',
    caption = 'By: Sam Burch | Data @Lahman'
  ) +
  geom_point(aes(size = IP, color = playerID, alpha = .8)) +
  geom_line(aes(color = playerID, alpha = .8), data = pitch_seasons |> filter(playerID == 'bumgama01'))

```

```
geom_line(aes(color = playerID, alpha = .8), data = pitch_seasons |> filter(playerID == 'riverma01'))
geom_line(aes(color = playerID, alpha = .8), data = pitch_seasons |> filter(playerID == 'kershcl01'))
geom_hline(yintercept = 1.30, color = 'darkgrey', linetype = 2)
```

WHIP Career Trajectories

Madison Bumgarner, Clayton Kershaw, and Mariano Rivera



By: Sam Burch | Data @Lahman

Note: The dotted grey line represents the approximate (modern) average WHIP in MLB – 1.3. Also, the lower the WHIP, the better.

Again, the brilliant play of these guys is clearly seen. Specifically, Kershaw and Rivera have shown elite peaks – with seasons dipping below .8 WHIP. For Bumgarner, his play (while still good) has been more up and down. On top of that, it is interesting how it took a few seasons for these pitchers to become really good. Maybe that suggests pitchers take longer to develop than hitters; however, these graphs don't necessarily prove that.

Question 3 Bob Gibson's 1968 season.

a)

```
Pitching |>
  filter(playerID == 'gibsobo01', yearID == '1968') |>
  select(GS, CG)
```

```
## # A tibble: 1 x 2
##   GS    CG
##   <int> <int>
## 1    34    28
```

28/34

```
## [1] 0.8235294
```

80% of the 34 games Gibson started (this season) were completed by him.

b)

```
Pitching |>
  filter(playerID == 'gibsobo01', yearID == '1968') |>
  select(BB, S0)
```

```
## # A tibble: 1 x 2
##       BB      S0
##   <int> <int>
## 1     62    268
```

268/62

```
## [1] 4.322581
```

About 4.3 strikeouts per walk.

c)

```
Pitching |>
  filter(playerID == 'gibsobo01', yearID == '1968') |>
  mutate(IP = IPouts / 3) |>
  select(IP)
```

```
## # A tibble: 1 x 1
##       IP
##   <dbl>
## 1  305.
```

Just over 300 IP.

d)

```
Pitching |>
  filter(playerID == 'gibsobo01', yearID == '1968') |>
  mutate(IP = IPouts / 3) |>
  mutate(WHIP = (BB + H) / IP) |>
  select(WHIP)
```

```
## # A tibble: 1 x 1
##       WHIP
##   <dbl>
## 1 0.853
```

About .85 WHIP.

Question 4 Jim Bunning's perfect game on Father's Day.

```
library(retrosheet)
```

```
## Warning: package 'retrosheet' was built under R version 4.2.2
```

```
##  
## For Retrosheet data obtained with this package:  
##  
## The information used here was obtained free of charge from  
## and is copyrighted by Retrosheet. Interested parties may  
## contact Retrosheet at "www.retrosheet.org"
```

```
retro_1964 = as_tibble(getRetrosheet(type = 'game', year = 1964))
```

a)

```
retro_1964 |>  
  filter(Date == '19640621',  
         VisTm == 'PHI',  
         HmTm == 'NYN',  
         DblHdr == '1') |>  
  select(Duration)
```

```
## # A tibble: 1 x 1  
##   Duration  
##   <int>  
## 1      139
```

The game was 2 hours and 19 minutes long.

b)

```
retro_1964 |>  
  filter(Date == '19640621',  
         VisTm == 'PHI',  
         HmTm == 'NYN') |>  
  select(Attendance)
```

```
## # A tibble: 2 x 1  
##   Attendance  
##   <int>  
## 1         0  
## 2      32026
```

The attendance value is equal to 0 because it is a double header. The attendance for the day is recorded for the second game that day, 32026.

c)


```
retro_1964 |>
  filter(Date == '19640621',
         VisTm == 'PHI',
         HmTm == 'NYN',
         DblHdr == '1') |>
  select(VisD, VisT, VisHR) |>
  sum()
```

```
## [1] 3
```

Three extra base hits by the Phillies.

d)

```
retro_1964 |>
  filter(Date == '19640621',
         VisTm == 'PHI',
         HmTm == 'NYN',
         DblHdr == '1') |>
  mutate(Vis_OBP = (VisH + VisBB + VisHBP) / (VisAB + VisBB + VisHBP + VisSF)) |>
  select(Vis_OBP)
```

```
## # A tibble: 1 x 1
##   Vis_OBP
##   <dbl>
## 1    0.333
```

.333 OBP by the Phillies.